

lab8b

September 2, 2024

```
[135]: import numpy as np# Summary of results
results = {
    'OLS': ols_mse,
    'Ridge': ridge_mse,
    'Lasso': lasso_mse
}

best_model = min(results, key=results.get)
print(f"The model with the highest accuracy (lowest MSE) is: {best_model}")

import pandas as pd
```

The model with the highest accuracy (lowest MSE) is: Lasso

```
[136]: dataset = pd.read_csv("Hitters.csv")
df = pd.DataFrame(dataset)
df
```

```
[136]:
```

	Unnamed: 0	AtBat	Hits	HmRun	Runs	RBI	Walks	Years	CAatBat	\
0	-Andy Allanson	293	66	1	30	29	14	1	293	
1	-Alan Ashby	315	81	7	24	38	39	14	3449	
2	-Alvin Davis	479	130	18	66	72	76	3	1624	
3	-Andre Dawson	496	141	20	65	78	37	11	5628	
4	-Andres Galarraga	321	87	10	39	42	30	2	396	
..	
317	-Willie McGee	497	127	7	65	48	37	5	2703	
318	-Willie Randolph	492	136	5	76	50	94	12	5511	
319	-Wayne Tolleson	475	126	3	61	43	52	6	1700	
320	-Willie Upshaw	573	144	9	85	60	78	8	3198	
321	-Willie Wilson	631	170	9	77	44	31	11	4908	

	CHits	...	CRuns	CRBI	CWalks	League	Division	PutOuts	Assists	\
0	66	...	30	29	14	A	E	446	33	
1	835	...	321	414	375	N	W	632	43	
2	457	...	224	266	263	A	W	880	82	
3	1575	...	828	838	354	N	E	200	11	
4	101	...	48	46	33	N	E	805	40	
..	

317	806	...	379	311	138	N	E	325	9
318	1511	...	897	451	875	A	E	313	381
319	433	...	217	93	146	A	W	37	113
320	857	...	470	420	332	A	E	1314	131
321	1457	...	775	357	249	A	W	408	4

	Errors	Salary	NewLeague
0	20	NaN	A
1	10	475.0	N
2	14	480.0	A
3	3	500.0	N
4	4	91.5	N
..
317	3	700.0	N
318	20	875.0	A
319	7	385.0	A
320	12	960.0	A
321	3	1000.0	A

[322 rows x 21 columns]

```
[137]: df.dtypes
```

```
[137]: Unnamed: 0      object
AtBat      int64
Hits       int64
HmRun      int64
Runs       int64
RBI        int64
Walks      int64
Years      int64
CAtBat     int64
CHits      int64
CHmRun     int64
CRuns      int64
CRBI       int64
CWalks     int64
League     object
Division   object
PutOuts    int64
Assists    int64
Errors     int64
Salary     float64
NewLeague  object
dtype: object
```

```
[138]: df.pop('Unnamed: 0')
```

```
[138]: 0      -Andy Allanson
        1      -Alan Ashby
        2      -Alvin Davis
        3      -Andre Dawson
        4      -Andres Galarrraga
        ...
        317    -Willie McGee
        318    -Willie Randolph
        319    -Wayne Tolleson
        320    -Willie Upshaw
        321    -Willie Wilson
Name: Unnamed: 0, Length: 322, dtype: object
```

```
[139]: df.head()
```

```
[139]:   AtBat  Hits  HmRun  Runs  RBI  Walks  Years  CAtBat  CHits  CHmRun  CRuns  \
0    293    66      1    30   29     14      1     293    66      1      30
1    315    81      7    24   38     39     14    3449   835     69     321
2    479   130     18    66   72     76      3    1624   457     63     224
3    496   141     20    65   78     37     11    5628  1575    225     828
4    321    87     10    39   42     30      2     396   101     12      48

      CRBI  CWalks  League  Division  PutOuts  Assists  Errors  Salary  NewLeague
0     29      14      A      E      446      33      20      NaN      A
1    414     375      N      W      632      43      10    475.0      N
2    266     263      A      W      880      82      14    480.0      A
3    838     354      N      E      200      11      3    500.0      N
4     46      33      N      E      805      40      4     91.5      N
```

```
[140]: df.describe()
```

```
[140]:   count    AtBat      Hits      HmRun      Runs      RBI      Walks  \
count  322.000000  322.000000  322.000000  322.000000  322.000000  322.000000
mean   380.928571  101.024845  10.770186   50.909938   48.027950   38.742236
std    153.404981   46.454741   8.709037   26.024095   26.166895   21.639327
min     16.000000    1.000000   0.000000   0.000000   0.000000   0.000000
25%    255.250000   64.000000   4.000000   30.250000   28.000000   22.000000
50%    379.500000   96.000000   8.000000   48.000000   44.000000   35.000000
75%    512.000000  137.000000  16.000000   69.000000   64.750000   53.000000
max    687.000000  238.000000  40.000000  130.000000  121.000000  105.000000

      Years      CAtBat      CHits      CHmRun      CRuns  \
count  322.000000   322.00000  322.000000  322.000000  322.000000
mean     7.444099  2648.68323  717.571429   69.490683  358.795031
std     4.926087  2324.20587  654.472627   86.266061  334.105886
min     1.000000   19.00000   4.000000   0.000000   1.000000
25%     4.000000   816.75000  209.000000  14.000000  100.250000
```

50%	6.000000	1928.00000	508.000000	37.500000	247.000000
75%	11.000000	3924.25000	1059.250000	90.000000	526.250000
max	24.000000	14053.00000	4256.000000	548.000000	2165.000000

	CRBI	CWalks	PutOuts	Assists	Errors \
count	322.000000	322.000000	322.000000	322.000000	322.000000
mean	330.118012	260.239130	288.937888	106.913043	8.040373
std	333.219617	267.058085	280.704614	136.854876	6.368359
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	88.750000	67.250000	109.250000	7.000000	3.000000
50%	220.500000	170.500000	212.000000	39.500000	6.000000
75%	426.250000	339.250000	325.000000	166.000000	11.000000
max	1659.000000	1566.000000	1378.000000	492.000000	32.000000

	Salary
count	263.000000
mean	535.925882
std	451.118681
min	67.500000
25%	190.000000
50%	425.000000
75%	750.000000
max	2460.000000

```
[141]: df.isnull().sum()
```

```
[141]: AtBat      0
Hits          0
HmRun         0
Runs          0
RBI           0
Walks         0
Years         0
CAtBat        0
CHits         0
CHmRun        0
CRuns         0
CRBI          0
CWalks        0
League        0
Division      0
PutOuts       0
Assists       0
Errors        0
Salary        59
NewLeague     0
dtype: int64
```

```
[142]: df['Salary'].fillna(value=df['Salary'].mean(),inplace=True)
```

```
[143]: from scipy import stats
import numpy as np

# Calculate Z-scores
z_scores = np.abs(stats.zscore(df.select_dtypes(include=np.number)))

threshold = 3
outliers = (z_scores > threshold)

print(np.where(outliers)[0])

df = df.drop(index=np.where(outliers)[0])
```

```
[ 30  32  73  73  80  82  82  84  96 100 112 113 114 121 136 163 163 179
 179 180 189 217 229 235 236 236 236 236 236 248 249 249 249 260 272
 274 276 278 292 302 302 302 302 302 306 310 313 315 320]
```

```
[144]: df.dtypes
```

```
[144]: AtBat          int64
Hits            int64
HmRun           int64
Runs            int64
RBI             int64
Walks           int64
Years           int64
CAtBat          int64
CHits           int64
CHmRun          int64
CRuns           int64
CRBI            int64
CWalks          int64
League          object
Division         object
PutOuts         int64
Assists         int64
Errors          int64
Salary          float64
NewLeague       object
dtype: object
```

```
[145]: # Convert categorical variables into dummy/indicator variables
df = pd.get_dummies(df, columns=['Division', 'League',
↪ 'NewLeague'],drop_first=True)
```

```
[146]: df
```

```
[146]:
```

	AtBat	Hits	HmRun	Runs	RBI	Walks	Years	CAtBat	CHits	CHmRun	\
0	293	66	1	30	29	14	1	293	66	1	
1	315	81	7	24	38	39	14	3449	835	69	
2	479	130	18	66	72	76	3	1624	457	63	
3	496	141	20	65	78	37	11	5628	1575	225	
4	321	87	10	39	42	30	2	396	101	12	
..	
316	221	53	2	21	23	22	8	1063	283	15	
317	497	127	7	65	48	37	5	2703	806	32	
318	492	136	5	76	50	94	12	5511	1511	39	
319	475	126	3	61	43	52	6	1700	433	7	
321	631	170	9	77	44	31	11	4908	1457	30	

	CRuns	CRBI	CWalks	PutOuts	Assists	Errors	Salary	Division_W	\
0	30	29	14	446	33	20	535.925882	False	
1	321	414	375	632	43	10	475.000000	True	
2	224	266	263	880	82	14	480.000000	True	
3	828	838	354	200	11	3	500.000000	False	
4	48	46	33	805	40	4	91.500000	False	
..	
316	107	124	106	325	58	6	535.925882	False	
317	379	311	138	325	9	3	700.000000	False	
318	897	451	875	313	381	20	875.000000	False	
319	217	93	146	37	113	7	385.000000	True	
321	775	357	249	408	4	3	1000.000000	True	

	League_N	NewLeague_N
0	False	False
1	True	True
2	False	False
3	True	True
4	True	True
..
316	True	True
317	True	True
318	False	False
319	False	False
321	False	False

```
[287 rows x 20 columns]
```

```
[147]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.metrics import mean_squared_error
```

```

# Prepare features and target variable
X = df.drop('Salary', axis=1) # Assuming 'Salary' is the target variable
y = df['Salary']

# Split data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)

# OLS Regression
ols_model = LinearRegression()
ols_model.fit(X_train, y_train)
ols_predictions = ols_model.predict(X_test)
ols_mse = mean_squared_error(y_test, ols_predictions)

# Ridge Regression
ridge_model = Ridge(alpha=1.0)
ridge_model.fit(X_train, y_train)
ridge_predictions = ridge_model.predict(X_test)
ridge_mse = mean_squared_error(y_test, ridge_predictions)

# Lasso Regression
lasso_model = Lasso(alpha=0.1)
lasso_model.fit(X_train, y_train)
lasso_predictions = lasso_model.predict(X_test)
lasso_mse = mean_squared_error(y_test, lasso_predictions)

print(f"OLS Mean Squared Error: {ols_mse}")
print(f"Ridge Mean Squared Error: {ridge_mse}")
print(f"Lasso Mean Squared Error: {lasso_mse}")

```

```

OLS Mean Squared Error: 59666.2269633068
Ridge Mean Squared Error: 59599.27851989654
Lasso Mean Squared Error: 59594.05119739496

```

```

/opt/anaconda3/lib/python3.11/site-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 5.386e+06, tolerance: 2.068e+03
    model = cd_fast.enet_coordinate_descent(

```

```

[148]: # Summary of results
results = {
    'OLS': ols_mse,
    'Ridge': ridge_mse,
    'Lasso': lasso_mse
}

```

```
best_model = min(results, key=results.get)
print(f"The model with the highest accuracy (lowest MSE) is: {best_model}")
```

The model with the highest accuracy (lowest MSE) is: Lasso

[]: